

# Introduction to Casual Inference - 097400

## Winter 2021 - Project Report

Asaf Gendler 301727715 Shalev Shaer 305252520

January 30, 2021

# **1 Introduction:**

Over the last 20 years, genome-wide association studies (GWAS) have established the association of thousands of genetic markers with phenotypes of interest, and these studies have the potential to shed light on essential scientific questions and to enable personalized medical care. As new methods for analyzing GWAS emerge, the problem that arises is that finding statistical association between genetic variants and some outcome, could also identify some irrelevant associations due to non genetic factors, and thus not necessarily indicate on a causal relation. There are methods to remove those irrelevant associations, however, they have no guarantee to ultimately remove them, thus the main goal over the next years will be to move from statistical association framework to causality framework, using the methods of this rising field. While the best way to identify causality is to perform a randomized controlled trial (RCT), this kind of trial is not possible when dealing with genotypes of humans. Instead, we can identify causality through observational study of father-mother-offspring trios data-sets due to the natural randomness of the recombination process of the offspring's genes. The paper [1] that we are basing on suggests a new analytic framework to identify causal effect of certain areas in the genome over some outcome through hypothesis testing using Conditional Randomized Tests(CRT). In this project we first review their methods and their mathematical justifications and then suggest a new approach to integrate randomness during training as well. We finish by showing some initial results that suggest that such an approach can indeed be useful.

## **2 Problem Formulation:**

### **2.1 Data-set Structure:**

To understand the problem, methods and the improvement we suggest, we first have to formulate the problem mathematically. To do so, we need to have a little bit of understanding of concepts in genetics, that will help us better understand our data-sets. The human cells have 36 DNA strands organized into 23 chromosomes pairs, In each chromosome pair, one strand is inherited from the father and one strand is inherited from the mother. Each one of these strands is called a haplotype. The haplotype consists of a chain of links called single-nucleotide polymorphisms or SNP. The reason they are called so is that they can appear in the population in one of two possible ways,

therefore they can be encoded as binary variables which can get the values of 0 or 1. See Figure 1 for visualization of the described above.




Figure 1: An illustration of haplotypes and SNPs. The chromosome pair consists of two DNA strands. Each strand is called an haplotype. Each single link in the haplotype is called a SNP and can take one of two forms in the population.

We consider a data-set of  $n$  father-mother-offspring trios, where for each one of them we have two haplotypes encodings of  $p$  SNP's values, all the haplotypes are of the same chromosome. The haplotypes of the father will be denoted as  $F^a$  and  $F^b$ , the mother's as  $M^a$  and  $M^b$ , and the offspring, which is our test subject as  $X^f$  and  $X^m$  as they are inherited from the father and

mother respectively. The data-set structure is thus as follows:

$$\left\{ \begin{array}{l} F^a = \begin{pmatrix} | & | & & | \\ F_1^a & F_2^a & \cdot & \cdot & \cdot & F_p^a \\ | & | & & | \end{pmatrix} \in \{0,1\}^{n \times p} \\ F^b = \begin{pmatrix} | & | & & | \\ F_1^b & F_2^b & \cdot & \cdot & \cdot & F_p^b \\ | & | & & | \end{pmatrix} \in \{0,1\}^{n \times p} \end{array} \right. \Rightarrow X^f = \begin{pmatrix} | & | & & | \\ X_1^f & X_2^f & \cdot & \cdot & \cdot & X_p^f \\ | & | & & | \end{pmatrix} \in \{0,1\}^{n \times p}$$

$$\left\{ \begin{array}{l} M^a = \begin{pmatrix} | & | & & | \\ M_1^a & M_2^a & \cdot & \cdot & \cdot & M_p^a \\ | & | & & | \end{pmatrix} \in \{0,1\}^{n \times p} \\ M^b = \begin{pmatrix} | & | & & | \\ M_1^b & M_2^b & \cdot & \cdot & \cdot & M_p^b \\ | & | & & | \end{pmatrix} \in \{0,1\}^{n \times p} \end{array} \right. \Rightarrow X^m = \begin{pmatrix} | & | & & | \\ X_1^m & X_2^m & \cdot & \cdot & \cdot & X_p^m \\ | & | & & | \end{pmatrix} \in \{0,1\}^{n \times p}$$

For easier notation we will denote the offsprings' matrix as  $X = X^f + X^m$ . The  $i$ -th row of this matrix  $X^{(i)}$  corresponds to the haplotypes of the  $i$ -th subject, and the  $j$ -th column  $X_j$  corresponds to the  $j$ -th SNP of all the subjects. We also denote  $A = \{F^a, F^b, M^a, M^b\}$  as the set of all parents haplotypes (A as a short for ancestors).

## 2.2 The goal

Our main goal is to establish causality in the Trio design, in a rigorous statistical sense. more precisely, we want to examine a causal connection between a SNP  $X_j$  or a group of SNPs  $X_C = \{X_j : j \in C\}$  to some phenomena or outcome  $Y$ . Differently from the main casual problems that we saw in the course, where the goal was to calculate the ATE or the ATT, our problem is only a binary problem that tries to answer the question whether or not there is a causal connection,

and doesn't try to estimate it's value. Mathematically we formulate the problem as a hypothesis testing problem as follows: given the examined SNP  $X_j$  and some outcome  $Y$  we want to test the null hypothesis:

$$H_0 : X_j \perp\!\!\!\perp Y \quad (1)$$

which means that this SNP and the outcome are statistically independent. Rejecting the null hypothesis will deduce that there is a statistical connection between the two, however, as we saw in the course statistical dependency and high correlation does not necessarily indicate a causal effect. For example, take  $Z$  to be some unobserved external variable that has a high casual effect on the outcome  $Y$ . Consider the case where the data-set contains two groups of populations, one group which tends to have the SNP  $X_j$  with a value of 1 and also tends to practice  $Z$ , and the other group which tends to have the SNP  $X_j$  with a value of 0 and also doesn't tend to practice  $Z$ . As a result, when testing the null hypothesis we suggested, we will probably reject it and find a statistical connection between  $X_j$  and  $Y$ , however, this relation was caused due the the correlation between  $X_j$  and  $Z$ , and the effect of  $Z$  on  $Y$ , and no due to a causal effect of  $X_j$  on  $Y$ .  $Z$  is in fact a hidden confounder as we saw in the course, and the causal graph will look like in Figure 2, which doesn't let us identify a causal effect of  $X_j$  on  $Y$ .




Figure 2: A possible causal graph. In this scenario we can't identify a causal connection between  $X_j$  and  $Y$  without observing  $Z$

As a result if we want to examine a causal effect, we need to do something smarter, as suggested in the next topic.

## 2.3 Identifying Causality

To solve the problem of some external unobserved hidden confounders, we need to address the recombination process of an offspring haplotypes from his parents haplotypes. Remember that the offspring has two haplotypes, one of them is a recombination of the two haplotypes of the father and the other is a recombination of the two haplotypes of the mother, an example for such a recombination can be seen in Figure 3 .




Figure 3: One offspring's haplotype is a recombination of the father's haplotypes and the other is a recombination of the mother's haplotypes

The crucial point for understanding, is that this recombination is a completely random process, with a known generation mechanism that can be simulated. Basically it means that given the haplotypes of the parents, we can stochastically model the generation process of the offspring haplotypes, simulate it, sample from it, etc. Mathematically it means that the distribution of  $X_j$  and actually all the SNPs is known given the parents haplotypes, which we denoted as the set  $A$ . From considerations of space, we won't fully detail the mathematical model of the recombination process, it is detailed in the paper and was implemented by us in python.

The main advantage of the randomness of the offspring's haplotypes recombination process is that if we go back to the scenario depicted in Figure 2, the causal graph should actually be depicted as in Figure 4.




Figure 4: The true casual graph,  $X_j$  is effected by  $Z$  only through  $A$  and thus  $A$  satisfies the backdoor criterion

What we see is that the external hidden confounder  $Z$  has a causal effect on  $X_j$  only through  $A$ , basically it means that any external hidden confounder can effect the parents haplotypes, but if we have observed them, then from there one the offspring haplotypes can be effected only by them and not by any unobserved hidden confounder. Understanding this, then instead of checking hypothesis Eq. (1), we will instead check the following null hypothesis:

$$H_0 : X_j \perp\!\!\!\perp Y | A \quad (2)$$

The main theorem of the paper shows that checking Eq. (2) is equivalent for checking the following:

$$H_0 : X_j \perp\!\!\!\perp Y | A, Z \quad (3)$$

where  $Z$  is any external unobserved variable. What we see here is basically the backdoor adjustment criterion which we saw in the course, the set  $A$  satisfies the backdoor criteria, corresponding to the causal graph in Figure 4, and thus conditioning on him is enough to identify causality between  $X_j$  and  $Y$ .

## 2.4 limitations

From what we have showed in section 2.3, we can deduce that if we have found a strong statistical connection between  $X_j$  and  $Y$  conditioned on  $A$ , which basically means that we will reject the null hypothesis Eq. (2), it has to mean that there is a causal effect between  $X_j$  and  $Y$ . However, this is not precisely the case. We have said that the recombination process of the offspring's

haplotypes is random conditioned on the parents, which is true, but the SNPs themselves for example,  $X_j$  and  $X_k$  are not independent on each other given the parents haplotypes  $A$ . This is due to the way the recombination occurs which we haven't discussed in details here but in short, the recombination tends to keep large chains of SNPs from the parents' haplotype, together in the offspring's haplotype. These dependencies can make the causal graph look like at Figure 5 for example.




Figure 5: SNPs can be dependent on each other given the parent haplotypes  $A$ , thus  $A$  itself doesn't satisfy the backdoor criterion anymore.

Analogously to our insight on the external counfounder  $Z$ , if we reject hypothesis Eq. (2), it could either mean that  $X_j$  indeed has an effect on  $Y$  or that some other SNP actually caused the effect. To deal with this issue, we will have to condition on all the other SNPs there we are not checking, and define a new null hypothesis:

$$H_0 : X_j \perp\!\!\!\perp Y | A, X_{-j} \quad (4)$$

where  $X_{-j}$  denotes all other SNPs besides  $X_j$ . Still, Hypothesis Eq. (2) that from now on we will call the global casual null hypothesis, has its benefits. This hypothesis testing still helps us find whether or not there is some SNP on the pair of haplotypes that might effect the outcome  $Y$ . If we reject the global null, it doesn't necessarily means that  $X_j$  has an effect on  $Y$  but it does mean that some SNP or a group of SNPs on the pair of haplotypes does have an effect. On the other hand, accepting the global null means not only that  $X_j$  doesn't have an effect but that actually all the SNPs on that haplotypes pair doesn't have an effect. This is why we call this hypothesis the global null, as it helps us to identify global causal effect between a pair of haplotypes and an outcome  $Y$ . If we want to further localize which SNP or group of SNPs actually caused the effect, we will need to use Eq. (4) which is therefore called the local casual null hypothesis. To check the local null we will have to know the distribution of  $X_j$  conditioned on his parents haplotypes,



and all the other SNPs. This is possible, and the paper details how to do so, however, because this makes the generation procedure much more complicated, and because our improvement can be cooperated with both the global null and the local null, we will deal with only the global null from now on.

We also remark that a few environmental factors of the parents can affect the inheritance process, such as the exposure of a parent to radiation, which changes the distribution of the offspring by increasing the frequency of mutations. This factors will of course break the backdoor criteria satisfied by  $A$ , however, those factors are very rare and can be neglectable most of the times.

### **3 Methods for solution:**

Now that we have understood which null hypothesis we want to check, let's discuss about the method to check it. We are going to check the global null hypothesis in the most generalized way, that is, examine the influence of a group of SNPs as follows:

$$\begin{aligned} H_0 : X_C \perp\!\!\!\perp Y | A \\ X_C = \{X_j : j \in C\} \end{aligned} \tag{5}$$

In order to do so, the paper suggests to use a version of the conditional randomization test suggested in [2] and calls it the Digital Twin Test (DTT). We will not proof the validity of the test but only explain how it works. This test is useful when you want to check a conditional hypothesis and you know how to sample from the conditional probability, which is exactly our case. We will describe the test in a more machine learning approach for better understanding of our improvement. First we take the data and split it into training and test sets. Then we fit some model between all the features, which are the SNPs in our case,  $X_C$  and  $X_{-C}$  to the outcome  $Y$ . We then calculate some error statistic  $T$  of the model on the test data, for example the MSE, and denote the value as  $t^*$ . Notice that this statistic  $T$  can in fact be any statistic we want (even doesn't have to integrate machine learning methods) and the validity of the test will still hold. Next, using our knowledge of the offspring's haplotypes given the parents haplotypes, we generate for each pair of parents in the test set  $K$  digital offsprings (They are like digital twins of the original offspring hence the name of the test), effectively creating  $K$  new test sets, and compute our statistic  $T$  on each one

of them, each result denoted as  $t_j$ . The critical point here is that when we create the digital offsprings, we only take the part  $\tilde{X}_C$  that we want to examine from it, and put in instead of the original  $X_C$  while leaving the other SNPs intact. At the end of the process we will have  $K$  values  $\{t_1, \dots, t_k\}$  and the original value  $t^*$ . All that is left to do is calculate the relative position of  $t^*$  among  $\{t_1, \dots, t_k\}$  values (for example(0.1)), this will be our p-value for the test (the test is rejected for any  $\alpha$  that is higher than the p-value). The reasoning behind this test, is that if the group  $X_C$  indeed has a statistical association with  $Y$  (and our model is well fitted), than by generating new digital offsprings with different values of  $X_C$ , the outcome values predicted by the model will change dramatically, hence the original error statistic will be relative small compared to the others and we will get a small p-value (which is s strong evidence for rejecting). On the contrary, if the value of  $X_C$  doesn't have any influence on the outcome, we expect the model to predict more or less the same outcome values just as before, hence  $t^*$  shouldn't be small compared to  $\{t_1, \dots, t_k\}$  and the p-value will be high, thus we will accept the null hypothesis. a summary of the algorithm is presented in Algorithem (1)

---

**Algorithm 1** The Digital Twin Test

---

- Input: data  $X$ , outcome  $Y$ , parental data  $A$ , region  $C$ , number of iterations  $K$

1. **Split** the data into train and test sets, and fit a model  $f(X_C, X_{-C})$  between the SNPs and the outcome on the train data.
2. **Compute** the error test statistic on the test data:

$$t^* = T(f(X_C, X_{-C}), Y)$$

3. **for**  $k = 1, \dots, K$  **Do**:

- (a) **Sample** the digital twins  $\tilde{X}$  based on the parents using the known  $P(X|A)$  and put  $\tilde{X}_C$  instead of the original  $X_C$ .
- (b) **Compute** the error test statistic using the digital twins:

$$t_k = T(f(\tilde{X}_C, X_{-C}), Y)$$

**end**

4. **Compute** the quantile of the true statistic  $t^*$  among the digital twin statistics  $\{t_1, \dots, t_k\}$

$$p_{val} = \frac{1 + \sum_{i=1}^K \mathbb{I}\{t_i \leq t^*\}}{K + 1}$$


---

## 4 Extension

### 4.1 general

Conditional Randomized Test (CRT) methods, such as the one we have described in the Digital Twin Test are now common in order to test hypothesis. CRT methods sample from a conditional distribution (in our context it was  $P(X_C|A)$ ) to create a "fake" sample  $\tilde{X}_C$ , and use the fake sample **only** in the test (as a counter sample). We suggest to involve the fake data  $\tilde{X}_C$  in the learning process as well (we can get infinite samples of  $\tilde{X}_C$  as we sample from a known distribution)

in order to increase the power of the test or to decrease the false detection rate. In this section we will combine  $\tilde{X}_C$  in a binary classifier, and examine whether the CRT give good results or not.

## 4.2 The proposed method

Our proposed method is as follow: First, we create the fake data  $\tilde{X}$  by sampling the desired region  $C$  from  $P(X|A)$  and substitute it with the desired region in  $X$ . Then we concatenate the outcome  $Y$  once with  $X$  and once with  $\tilde{X}$ , thus creating real data and fake data. We add labels to each one of this data sets, the real data gets the label "1" and the fake data the label "0". We then train a binary classifier with the both the real and fake data that tries to extinguish who is real and who is not. The  $T$  statistic in the CRT, which liked said before can be any statistic, will now be the sum of the probabilities of getting "1" in the output of the classifier. The full algorithm is described below.

---

**Algorithm 2** The Digital Twin Test with binary classifier
 

---

- Input: true data  $X$ , outcome  $Y$ , parental data  $A$ , binary classifier, region  $C$ , number of iterations  $K$

1. **Create**  $\tilde{X}$  by sampling the digital twins like in DTT  $\tilde{X} = (\tilde{X}_C, X_{-C})$
2. **Create** the training data by concatenating  $Y$  to  $X$  and to  $\tilde{X}$ , and add the label "1" to  $X$  and "0" to  $\tilde{X}$
3. **Fit** the classifier with both  $(X, Y, 1)$  and  $(\tilde{X}, Y, 0)$  as one training set
4. **Compute** the test statistic on the test data:

$$t^* = - \sum_{i \in Test} P(classifier(x_i, y_i) = 1)$$

5. **for**  $k = 1, \dots, K$  **Do**:

- (a) **Sample** the digital twins  $\tilde{X}$  based on the parents using the known  $P(X|A)$  and put  $\tilde{X}_C$  instead of the original  $X_C$ .
- (b) **Compute** the test statistic using the digital twins:

$$t_k = - \sum_{i \in Test} P(classifier(\tilde{x}_i, y_i) = 1)$$

**end**

6. **Compute** the quantile of the true statistic  $t^*$  among the digital twin statistics  $\{t_1, \dots, t_k\}$

$$p_{val} = \frac{1 + \sum_{i=1}^K \mathbb{I}\{t_i \leq t^*\}}{K + 1}$$


---

The reasoning behind our new proposal is as follows. If the set of SNPs  $X_C$  has a large influence on the outcome  $Y$ , then we expect our classifier to train well, and be able to distinguish well between real examples and the fake ones. In this case, when we will generate the digital twins on the test set, the classifier will tend to give higher probabilities to the original data over the generated one.

On the contrary, if the SNPs doesn't influence  $Y$  at all, than our trained classifier will act poorly, not able to distinguishing between real and fake, hence the probabilities it will give to the digital twins during test will be more or less the same to the probabilities on the original set.

## 5 Simulations and Results:

First, let us describe the settings. We create a synthetic population of  $n = 250$  parent-child trios and generate an outcome coming from a sparse regression model  $Y^{(i)} = \beta^T X^{(i)} + v$  where  $v \sim N(0, 1)$ . We use  $p = 100, 150$  SNPs from a chromosome with a width of 50 Mb, and  $\beta$  is sparse and has only 10 non zeros entries of equal strength. We used four different signal strength  $c = 1, 2, 3, 4$ , and for every value of  $c$  we used the suggested algorithm 20 times. We choose Quadratic Discriminant Analysis (QDA) as our classifier. These experiments demonstrate a situation of a causal region, therefore we expect low p-values. We plot a plot-box for the extracted p-values for every setting as can be seen in Figure 6 and Figure 7. As we can see, our method behave as expected - the higher the signal strength the smaller the p-value and the narrower the standard deviation of the p-values. Furthermore, the more SNPs there are (the greater the  $p$ ), the higher the p-values and wider the standard deviation. The original DTT on this setting always results in a minimal close to zero p-value so we didn't plot it as it does not bring more information than this mention. For a final experiment we produce the outcome  $Y$  from a random distribution  $N(0, 1)$ , and ran the suggested algorithm with QDA classifier. The results are in Figure 8. We can see that for a noisy outcome, unrelated to the SNPs, the p-values are high and with high standard deviation, as expected.

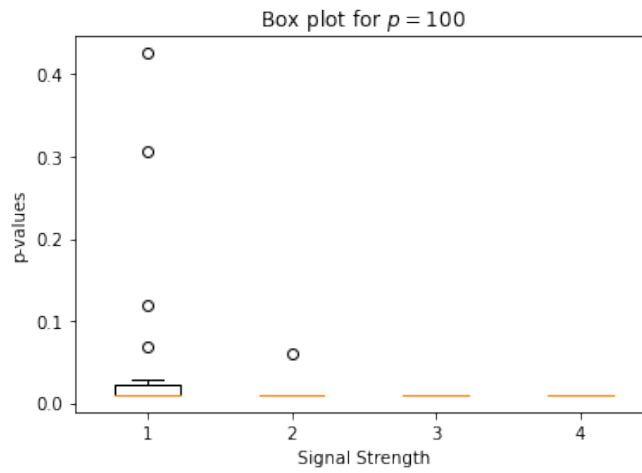


Figure 6: A box plot for  $p = 100$  and  $k = 10$

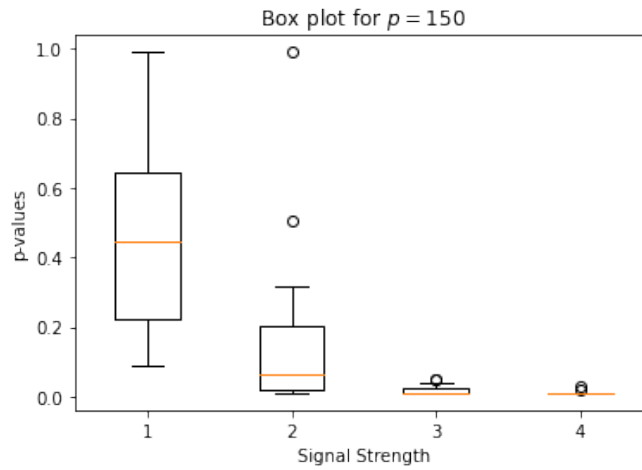


Figure 7: A box plot for  $p = 150$  and  $k = 10$

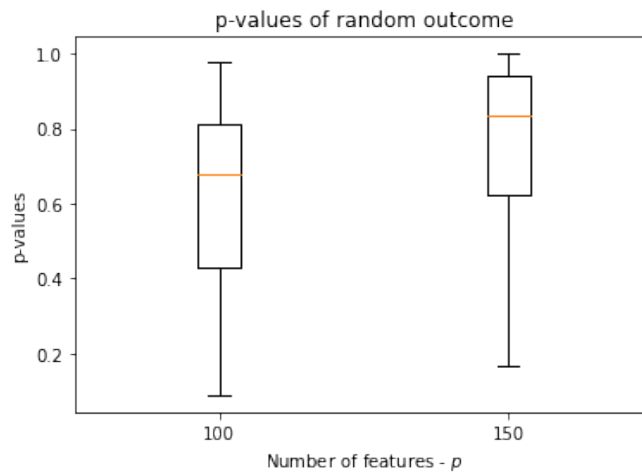


Figure 8: A box plot for  $p = 100$ ,  $p = 150$  and a noisy output  $Y$  unrelated to the SNPs

## 6 Conclusion and Future Work:

Our method does not outperform the DTT as described above, but it does work well and behave as expected. The main goal of this project was to examine whether our idea to use  $\tilde{X}$  in the training process make sense or not, and it seems that it does. For future work one can consider involve  $\tilde{X}$  in more complex methods (e.g Neural Networks) and check whether it increases the power and \ or decrease the false detection rate.

It is important to note that although we didn't test our method with the local causal null hypothesis, we expect similar result as the framework is the same and only the conditional sampling procedure changes. Another aspect that both us and the paper didn't investigate is the use of multiple hypothesis testing methods, such as FDR, to unite examining of different SNPs regions into one more powerful test. Of course that our method can be extended to this filed as well, as it is generic.

## References

- [1] S. Bates, M. Sesia, C. Sabatti, and E. Candès, “Causal inference in genetic trio studies,” *Proceedings of the National Academy of Sciences*, vol. 117, no. 39, p. 24117–24126, Sep 2020. [Online]. Available: <http://dx.doi.org/10.1073/pnas.2007743117>
- [2] E. Candès, Y. Fan, L. Janson, and J. Lv, “Panning for gold: “model-x” knockoffs for high dimensional controlled variable selection,” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, vol. 80, no. 3, p. 551–577, Jan 2018. [Online]. Available: <http://dx.doi.org/10.1111/RSSB.12265>